



# Forgejo with SSO

Virtualization and Cloud Computing  
a.a. 2023/24

Authors:

Gabriele Dellepere – S4944557

Lorenzo Foschi – S4989646

Kevin Cattaneo – S4944382





# Main challenges

01

Docker

02

NFS

03

Swarm

04

Registry

05

Traefik

06

PostgreSQL

07

Keycloak

08

Forgejo

09

Grafana

10

Prometheus

11

Loki

12

Promtail

...yes, **everything**






# Oh I forgot... Vagrant



## Lacks of pre-requisites

*vagrant plugin install*  
**vagrant-sshfs** command  
was needed before *vagrant*  
*up*



## Tweaking hardware configuration

In order to start the giant we  
needed to set the right  
combination of RAM  
allocated to each VM





# Vagrant bugs



#1

## Default provider

Solution was to set  
**DEFAULT\_PROVIDER** =  
'vmware' instead of  
**DEFAULT\_PROVIDER** =  
'libvirt' in Vagrantfile



#2

## AudioAdapter

The box.ovf created by the  
VagrantFile crashes when  
Linux + Virtual Box are used on  
the incompatibility of the  
**AudioAdapter** (by default it is  
used the one of Windows)






# Docker, Swarm, NFS




## Refining the solution

There was a lot of work on trying to refine multiple times the solution while understanding the **Ansible syntax**, since these were the first touches



## Diving in scalable solution

Some effort was involved when we delved into purple task: scalable solution, trying to understand who **loop** on what variable



## Is really exported in that way?

While developing NFS tasks, we questioned about the workflow creation, export, update, and the necessity of **fstab**





# Docker Registry #1



## Pre-tasks

We needed to install **apache-utils** for `htpasswd` and also dependencies like **passlib** to make the tasks work



## Permissions

We also spent some time searching for the optimal permission, in particular for the '0400' of the **RSA private key** (only readable by owner and root)





## Docker Registry #2



### TLS

To enable TLS we added some environment variables that allowed us to refer to certificate and key, but we also needed to specify **tls: true** in the Run Registry task



### Backfire of our work

Some fixes arrived after developing the following tasks, for example the **creation of CSR before** the creation of the self-signed certificate





# Swarm Services: Database / PostgreSQL




## Which entrypoints?

One database for each service, it was clear; but a lot of time was wasted in order to understand we were required to touch the .sql and to create volumes



## The dilemma of Postgre 'role'

By looking to the logs we understood and loved the fact that Postgre uses different syntax for **role** and **user** creation



## The elephant logo has a meaning

As we developed the following tasks we noticed the **slowness** of Postgres





# Swarm Services: Traefik #1



## The right certificate...

By looking to the logs we discover a bug on the format of the certificate: we needed a **.pem** instead of **.cert**



## ...in the right directory

We also added **/etc/ssl/traefik** to create the proper directory for placing and locate the newly created certificates



## Swarm Services: Traefik #2



### Overlay network

In order to allow the services to talk each other by passing through traefik, we put them in an **overlay\_net**




### The hosts file

In order to make curl retrieving correctly data and make https services reachable, we wrote the right host (\*.vcc.local) in a file, so that we can mount the **/etc/hosts** volume in the services that need it



# Swarm Services: Keycloak #1



## Indentation is important

Mostly when you talk about YAML, infact we had Keycloak db failure since our labels are referred to compose and not the swarm service



## Removal of transactions

When we implemented the .sql files we also used **transactions** but this created some conflicts so we removed them



## Swarm Services: Keycloak #2




### Postgre permissions

We discovered that some permissions are not given to a user unless he is the **creator** of that db; so we scheduled first the creation of user and then of the db



### Why so shy?

Keycloak was not visible on `auth.vcc.local`, that was caused by the unknown port to be contacted to access the service from the load balancer; so we specified it with:  
**`traefik.http.services.auth.loadbalancer.server.port=8080`**



### Creation of realm

The creation of the realm was one of the most difficult task: we then succeeded only when the problem occurred in the successive tasks. Until that our SSO was not working



# Swarm Services: Forgejo



## Update ca-certificates

Was an open debate until we found out that Forgejo is an alpine distribution and so we wrote the right way of installing the command to update the certificates



## Forcing the wait

We had also some db problems related to the slowness of Postgres, so we **forced some sleep** during the Forgejo process



## Alpine container

Just during developing phase we also created an alpine container to better understand the default package installer of the distribution and also if the services were reachable



# Swarm Services: Grafana #1



## Unclear documentation

Online there is **unclear documentation** about Grafana, so trying to master the interface has been challenging



## OAuth

The major difficulty was enabling the SSO on Grafana and the whole problem can be translated to choosing the **right environment variables**



## Other OAuth problems

- For Grafana: missing the line **export GF\_SERVER\_ROOT\_URL=https://mon.vcc.local**
- For Forgejo: the **--name** option at the end of entrypoint.sh calls was missing



## Swarm Services: Grafana #2



### Journal the container

Due to a bit unclear documentation we couldn't manage to start the container: we used **journalctl -u docker.service** founding out that we were mounting the wrong volume



### Dashboards

Dashboards took a lot of time, since they need to show the monitoring. Also it was all done by hand because **no template was suitable**



## Swarm Services: Grafana #3



### SSO Grafana as a real admin

To access with Keycloak's SSO in Grafana as Admin users (so being able to create dashboards accessing with SSO) we had to retry a lot of realms. Then we found the right realm configuration and the right '**attribute role**' string syntax.



### **!= error let us discover that...**

We weren't selecting the proper UIDs for data sources (Loki and Prometheus) in dashboard json's, linking with the manually created ones in the data sources YAML





# Swarm Services: Prometheus



## OAuth reverse proxy

The major difficulty has been how to create the reverse proxy to allow Prometheus to be accessed with Keycloak



## Scraping registry metrics

Accessing metrics from registry was hard. We added **dockerswarm\_sd\_configs** at the end of prometheus.yml, with a **port:5051** replacement (the one that we opened in task 12 as DEBUG PORT for metrics)





# Swarm Services: Promtail, Loki, Fluent-bit



## State your purpose

Overall they have not been so difficult to develop apart the understanding about which was the **purpose** of each service in our scenario



## Loki's default UID/GID

Loki uses as default UID/GID = 10001 that lacks of mounting permissions. We decided not to leverage the user to root, so we specify the mounting with the **right directories** while having the permissions






## Some more details

We can also confirm that:

- Our system works when we **scale up** to  $N\_TARGET > 2$
- We committed all the suggestion provided by **Ansible Lint**, so our scripts follow the good-programming manners for Ansible
- Our privacy-critical files have been **ansible-vaulted**



# Time spent for the project



## First month

December 2023 – 80%

## Second month

January 2024 – 20%

**18 days**



spent from the 20 December to 7 January

# Thanks!

Credits for presentation template to **Slidesgo** and  
**Freepik** for infographics

