

Assignment 1 - BASC

BINARY ANALYSIS AND SECURE CODING - UniGE

Kevin Cattaneo - S4944382

Index

Index	1
Extract	1
ELF 1	2
Solution.....	2
ELF 2	4
Solution.....	4
ELF 3	5
Solution.....	5
ELF 4	6
Solution.....	7
Extra	8

Extract

The goal of this assignment is to get acquainted with the ELF file format. Unzip `ELF_files.zip`, and you'll have a bunch of files to work with. All of them hide one flag, that is, a string of the form `BASC{...}`, where `"..."` is at least eight characters long. In particular, `BASC{3T0N5}`, which is something you may come across, is not the intended flag for the corresponding file.

Given that, I will `unzip ELF_files.zip`.

ELF 1

First of all I run a bunch of commands to make sure of what I'm handling.

```
file elf1
elf1: ELF 64-bit LSB pie executable, x86-64, version 1 (SYSV),
dynamically linked, interpreter /lib64/ld-linux-x86-64.so.2,
BuildID[sha1]=cb565ed38996d33e5c6966bbe1699175b9924c75, for
GNU/Linux 3.2.0, stripped
>> It seems like a plain x86-64 ELF executable.
```

```
./elf1
Nothing to see here...
```

```
strings elf1 - no flags found
```

Then I explore the binary with `hte elf1` and discover the `BASC{3TON5}`, that is not intended to be a valid flag.

Solution

Run the `readelf` command, just after a very big look to the man. The flag can be found with different command-flag combinations.

```
readelf -a elf1 - shows everything that will follow
```

```
readelf -S elf1
```

There are 43 section headers, starting at offset 0x3158:

Section Headers:

[Nr]	Name	Type	Address	Offset
	Size	EntSize	Flags Link Info Align	
[17]	B	PROGBITS	00000000000011e5	
	000011e5			
	000000000000000b	0000000000000000	AX 0 0	1
[18]	A	PROGBITS	00000000000011f0	
	000011f0			
	000000000000000b	0000000000000000	AX 0 0	1
[19]	S	PROGBITS	00000000000011fb	
	000011fb			
	000000000000000b	0000000000000000	AX 0 0	1
[20]	C	PROGBITS	0000000000001206	
	00001206			
	000000000000000b	0000000000000000	AX 0 0	1
[21]	{	PROGBITS	0000000000001211	
	00001211			
	000000000000000b	0000000000000000	AX 0 0	1

[22] s	PROGBITS	0000000000000121c		
0000121c				
0000000000000000b	0000000000000000	AX	0	0 1
[23] 3	PROGBITS	00000000000001227		
00001227				
0000000000000000b	0000000000000000	AX	0	0 1
[24] c	PROGBITS	00000000000001232		
00001232				
0000000000000000b	0000000000000000	AX	0	0 1
[25] T	PROGBITS	0000000000000123d		
0000123d				
0000000000000000b	0000000000000000	AX	0	0 1
[26] i	PROGBITS	00000000000001248		
00001248				
0000000000000000b	0000000000000000	AX	0	0 1
[27] 0	PROGBITS	00000000000001253		
00001253				
0000000000000000b	0000000000000000	AX	0	0 1
[28] N	PROGBITS	0000000000000125e		
0000125e				
0000000000000000b	0000000000000000	AX	0	0 1
[29] 5	PROGBITS	00000000000001269		
00001269				
0000000000000000b	0000000000000000	AX	0	0 1
[30] }	PROGBITS	00000000000001274		
00001274				
0000000000000000b	0000000000000000	AX	0	0 1
[31] .fini	PROGBITS	00000000000001280		
00001280				
0000000000000000d	0000000000000000	AX	0	0 4v

readelf -l elf1:

Section to Segment mapping:

Segment	Sections...
00	
01	.interp
02	.interp .note.gnu.property .note.gnu.build-id .note.ABI-tag .gnu.hash .dynsym .dynstr .gnu.version .gnu.version_r .rela.dyn .rela.plt
03	.init .plt .plt.got .plt.sec .text B A S C { s 3 c T i 0 N 5 } .fini
04	.rodata .eh_frame_hdr .eh_frame
05	.init_array .fini_array .dynamic .got .data .bss
06	.dynamic
07	.note.gnu.property
08	.note.gnu.build-id .note.ABI-tag
09	.note.gnu.property

```
10      .eh_frame_hdr
11
12      .init_array .fini_array .dynamic .got
```

Multiple sections at segment 03 have been created with their headers (so that we can see them from both Section and Program Headers), that compose the 8-length content of the flag: **BASC{s3cTi0N5}**, thanks for the README hint.

After a little investigation, the fact that BASC{3T0N5} is a subset of the real flag is because of the section ".shstrtab", that is a section that maps other sections and re-utilize bytes when certain elements (e.g. letters) are double used; indeed if we look at the offsets pointed by the content of this section we can also find the missing letters.

ELF 2

```
file elf2
elf2: ELF 32-bit LSB executable, ARM, EABI5 version 1 (SYSV),
statically linked, for GNU/Linux 2.6.26,
BuildID[sha1]=84f38995563fddf7e782f901ad6c84a63c7e01f9, stripped
>> It seems to be like an ARM executable
```

```
./elf2
zsh: exec format error: ./elf2
```

```
strings elf2 - no flag found
```

In the first place I tried to edit with hte editor to change some values in the ELF header to resemble a x86_64 and also tried to change endianness, but with no result.

With further exploration I verified that it is indeed an ARM, since we have some legit ARM sections in it.

```
readelf -l elf2
Section to Segment mapping:
Segment Sections...
 00      .ARM.exidx
 01      .note.ABI-tag .note.gnu.build-id .init .text
__libc_freeres_fn .fini .rodata __libc_subfreeres __libc_atexit
.ARM.extab .ARM.exidx .eh_frame
 02      .tdata .init_array .fini_array .jcr .data.rel.ro .got
.data .bss __libc_freeres_ptrs
 03      .note.ABI-tag .note.gnu.build-id
 04      .tdata .tbss
 05
```

Solution

After reviewing the given material, I installed the qemu-user-binfmt virtualizer to run ARM. Just launch elf2 again:

```
./elf2
BASC{ARMed_&_d4ng3r0uS}
```

ELF 3

```
file elf3
elf3: data
```

>> This differs a lot from the other ELF files, in some sense it cannot recognize it is an ELF at all.


```
./elf3
zsh: exec format error: ./elf3
>> As expected it will not run
```

```
strings elf3 - no flag found
```

But I discovered from this string `"/lib64/ld-linux-x86-64.so.2"`, that the file I'm working with may be some kind of ELF 64-bit executable for x86_64 and not a generic "data" file.

Solution

After viewing the file with `hxd`, I discover that the first identifier bytes for the magic number of the ELF type are different: the letters "elf" are in lowercase.



By changing them to the correct ones: `7f 45 4c 46` I get a correct ELF executable.

```
./elf3
BASC{cAs3_maTT3rS}
```

Another way is by running:

```
readelf -l elf3
readelf: Error: Not an ELF file - it has the wrong magic bytes at
the start
```

ELF 4

```
file elf4
```

```
elf4: ELF 64-bit LSB executable, x86-64, version 1 (SYSV),
statically linked,
BuildID[sha1]=b7c34e2c255427db1a4bf366a804fc3d67ff35f0, for
GNU/Linux 3.2.0, stripped
```

>> It seems a plain x86-64 ELF 64-bit executable

```
./elf4
```

```
[1] 6588 segmentation fault (core dumped) ./elf4
```

strings elf4 – there is the universe made into strings in this program, but no flag found

I tried to edit with hte editor some ELF header bytes with different combinations, but the ELF always tells about a segmentation fault or about an exec format error after those edits.

After a lot of trial and error, I investigated with gdb debugger and I see the program stops immediately with at the very start (with both start and starti) at address:

```
0x401c60 endbr64.
```

Because of:

Program terminated with signal SIGSEGV, Segmentation fault.

This address is in the .text section, which holds the code of the program.

```
[*] /home/keat/Desktop/elf-headers/elf
* ELF section headers at offset 0x000c2788
[+] section 0:
[+] section 1: .note.gnu.property
[+] section 2: .note.gnu.build-id
[+] section 3: .note.ABI-tag
[+] section 4: .rela.plt
[+] section 5: .init
[+] section 6: .plt
[-] section 7: .text
  name string index
  type              0000004f
  flags             00000001 (progbits)
  address           0000000000000006 details
  offset            00000000004011a0
  size              0000000000011a0
  link              0000000000091820
  info              0000000000000000
  alignment         00000000
  entsize           00000000
[+] section 8: __libc_freeres_fn
[+] section 9: .fini
[-] section 10: .rodata
```

I check for details of .text, to make clear that everything is as it should be, in particular about correctness of permissions, specifically if it is executable. It is indeed.

```
[*]
[00] writable      0
[01] alloc         1
[02] executable    1
[03] ???           0
[04] merge         0
[05] strings       0
[06] info link     0
[07] link order    0
[08] OS non-conforming 0
```

Solution

After checking the `.text` section I also tried to look at the program header table to see also details about the segment that contains the `.text` section.

```
readelf -l elf4
```

Section to Segment mapping:

```
Segment Sections...
 00      .note.gnu.property .note.gnu.build-id .note.ABI-tag
.rela.plt
 01      .init .plt .text __libc_freeres_fn .fini
 02      .rodata .stapsdt.base .eh_frame .gcc_except_table
 03      .tdata .init_array .fini_array .data.rel.ro .got
.got.plt .data __libc_subfreeres __libc_IO_vtables __libc_atexit
.bss __libc_freeres_ptrs
 04      .note.gnu.property
 05      .note.gnu.build-id .note.ABI-tag
 06      .tdata .tbss
 07      .note.gnu.property
 08
 09      .tdata .init_array .fini_array .data.rel.ro .got
```

I see that the `.text` section is in the segment 01, that is this segment:

Type	Offset	VirtAddr	PhysAddr
	FileSiz	MemSiz	Flags
Align			
LOAD	0x00000000000001000	0x00000000000401000	
	0x00000000000401000		
		0x000000000009366d	0x000000000009366d R
	0x1000		

To verify the correctness I compared that segment with the one correspondent (so the one that contains `.text`) in a plain x86_64 ELF executable (a simple `printf` in a `main`) of mine: I observed that the `elf4` is missing the `E`(xecutable) permission on the segment above. This is also visible from `hte`.

```
* ELF program headers at offset 0x00000040
[+] entry 0 (load)
[-] entry 1 (load)
  type                                00000001 (load)
  flags                               00000004 details
  offset                             00000000000001000
  virtual addr [ * ]
  physical add [00] executable      0
  in file size [01] writable        0
  in memory si [02] readable        1
  alignment
[+] entry 2 (l
[+] entry 3 (l
```

So by editing this permission flag executable to 1 I repaired `elf4`.

```
./elf4
```

```
BASC{no_eXec_no_party}
```


Extra

"Flags are not four, but every file has one flag".

If you check:

```
file ELF_files.zip
```

```
ELF_files.zip: ELF 64-bit LSB pie executable, x86-64, version 1
(SYSV), dynamically linked, interpreter
/lib64/ld-linux-x86-64.so.2,
BuildID[sha1]=69e1fd0c892100efe04b1cdb0628433d7bd26038, for
GNU/Linux 3.2.0, stripped
```

>> But wait, it should have been just a ZIP! So it is a polyglot file.

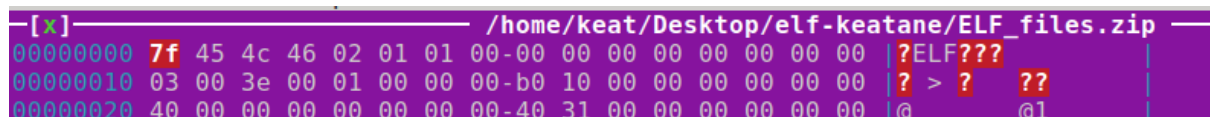
So if I run it:

```
chmod +x ELF_files.zip
```

```
./ELF_files.zip
```

```
BASC{can_U_run_a_ZIP?}
```

Indeed, if we explore the first part of the ZIP files, it has the ELF identifier.



```
[*] /home/keat/Desktop/elf-keatane/ELF_files.zip
00000000 7f 45 4c 46 02 01 01 00-00 00 00 00 00 00 00 00 00 |?ELF???
00000010 03 00 3e 00 01 00 00 00-b0 10 00 00 00 00 00 00 00 |? > ? ??
00000020 40 00 00 00 00 00 00 00-40 31 00 00 00 00 00 00 |@ @1
```

The doubt that can arise is how: as we have seen file extensions are immaterial and, as course notes says:

“Different parsers can interpret the same sequence differently

- ZIP parsers look for the “End of Central Directory” from the **end** of the file
- ELF parsers expects an header at the **beginning**”